

Non-Docker deployment - prior requirements

Non-Docker deployment - prior requirements

Install these **before** running `./install.sh`. The installer can install some packages on Debian/Ubuntu via `apt-get`;

Required (all deployments)

Debian/Ubuntu one-liner:

```
sudo apt-get update && sudo apt-get install -y \
  build-essential cmake curl python3 python3-pip python3-venv git
```

Required for 150?300 user production (non-Docker)

```
sudo apt-get install -y nginx
```

Use `./scripts/run_native_300.sh` after install - it sets worker counts and starts nginx + gunicorn.

Optional - GPU backends

```
ENABLE_CUDA=1 ./install.sh
```

```
ENABLE_OPENCL=1 ./install.sh
```

Optional - testing & profiling

```
sudo apt-get install -y valgrind
```

Scaling tiers (non-Docker)

Environment (written by `install.sh` to `.env.nanoserve`):

```
export NANOSERVE_NUM_WORKERS="$(nproc)"    # overridden by run_native_300.sh
```

```
export NANOSERVE_MAX_BATCH="32"
```

```
export NANOSERVE_MAX_QUEUE="512"
```

```
export LD_LIBRARY_PATH="$PWD/allocator/target/release:$LD_LIBRARY_PATH"
```

```
export NANOSERVE_ENGINE_LIB="$PWD/engine/build/libnanoserve_engine.so"
```

```
export PYTHONPATH="$PWD:$PYTHONPATH"
```

Network

- Port **8000** - HTTP API + Web UI (nginx in production mode)

- Ports **8001?8004** - internal gunicorn (production only)

What the installer does not provide

- Docker / container runtime

- NVIDIA Container Toolkit (Docker GPU only)

- Systemd unit files (add your own for production)

- TLS certificates (terminate TLS at nginx or a reverse proxy)

See SETUP.md and USAGE.md for step-by-step setup and daily use.